

Chapter 9: Virtual Memory Management

1. Background

- **Virtual memory**

: separation of user **logical memory** from **physical memory**



- allows **only part of the program** to be in memory for execution
- allows logical address space to be **much larger** than physical address space
- allows address spaces to be **shared** by several processes
- allows for **more efficient process creation**

• Virtual memory can be implemented via:

- **Demand paging**
- **Demand segmentation**

2. Virtual Memory : larger than physical memory

- Virtual memory involves the separation of **logical** memory and the **physical** memory
- The separation allows an extremely **large** virtual memory for programmers.
- The programmer no longer need to worry about the amount of physical memory available.
- **MMU maps virtual address to the physical address.**




3. Virtual-address Space

- Virtual address space
 - The **logical view of how process is stored in memory**
 - A process begins at
 - a certain logical address
 - address 0
 - A process exists in **contiguous memory**
 - A process consists of **4 segments**
 - **Code, Data, Heap, Stack**
 - **Heap** grows upward in memory
 - **Stack** grows downward in memory
- The **large blank space** between the heap and the stack require actual physical pages only if the heap or stack grows.

4. Shared Library using Virtual Memory

- Virtual memory allows **files and memory** to be **shared** by a number of processes through **page sharing**.
 - allows system libraries to be shared by several processes.
 - allows processes to create a shared memory
 - allows pages to be shared during process creation with fork()




5. Demand Paging

- is a **virtual memory management strategy**
- **Pages are loaded into memory only when they are needed during program execution** → **dynamic loading**
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed → reference to it
 - invalid reference → abort
 - not-in-memory → bring to memory

6. Demand-paging system with swapping

- A **demand-paging system** is similar to a **paging system with swapping** where processes reside in secondary memory (disk)



- A process address space is a sequence of pages in demand paging
- When a process is created
- **Lazy swapper (pager)** is used
 - Never swaps a page into memory unless that page will be needed.
- **pure demand paging**

* Demand paging은 swapping을 한다는 점에서 paging과 유사

* 차이: swapping할 때, 프로세스를 실행하기 위해 프로세스의 모든 페이지가 아닌, 필요한 페이지만 메모리로 불러들인다

* Swapper는 전체 프로세스를 다루지만 / Pager는 프로세스의 개별 페이지와 관련

7. Valid-Invalid Bit


- In demand-paging system, some pages reside in the memory and others in the disk
 - needs a H/W support to distinguish it
- a **valid-invalid bit** is associated with each page table entry (**1 → in-memory, 0 → not-in-memory**)
- Initially valid – invalid bit is set to **0** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	1
	1
	0
	1
	0

page table


- During **address translation**, if **valid-invalid bit** in page table entry is 0 → **page fault**

8. Page Table when some pages aren't in main memory



- When a page is loaded in memory
 - **valid-invalid bit** is **set :1**
- When a page is in disk
 - “ **unset :0** ”
- The **frame value** contains the **address of the page on the disk**
- **Page Fault**
 - When a process tries to access a page not in memory (주소 변환을 하는 동안 페이지 테이블의 해당 엔트리의 valid-invalid bit가 0이면)
 - **Page fault trap** occurs.

9. Steps in Handling a Page Fault



1. Check the reference is valid or not;
If not, terminate the process.

1. Check the page is in memory with valid-invalid bit in page table;
→ If not,
2. Page fault to OS
3. Find the page on the disk
4. Find a free frame
4. Read the page into the newly allocated frame from disk
5. Update the page table
6. Restart the instruction

- Transformed into physical address
- Access physical memory to get data

* Page Fault Service Routine

1. 참조한 주소가 유효한 접근 (valid access)인지 결정하기 위해 내부 테이블(PCB)을 확인한다.
2. 유효하지 않은 접근이면 종료하고, 유효하나 page fault이면 page in한다.
3. 비어있는 프레임을 찾는다.
4. 디스크상의 원하는 페이지의 위치를 찾아서 이를 비어있는 프레임으로 읽어드린다.
5. 해당 페이지에 대한 페이지 테이블의 엔트리를 수정한다.
6. 사용자 프로세스를 다시 시작한다.

10. What happens if there is no free frame?

- **Page Fault** → find a target frame for loading a new page
- What happens if there is **no free frame**?
- One of the valid frame need to be replaced with the new one
- **Page replacement** – find a frame in memory, but not really in use, swap it out, swap a new page in
- Several algorithms for page replacement
- Performance – want an algorithm which will result in **minimum number of page faults** with a given reference string

11. Performance of Demand Paging

- **Page Fault Rate** $0 \leq p \leq 1.0$
- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault
 - **Effective Access Time (EAT)**
$$EAT = (1 - p) \times [\text{memory access time}] + p \times [\text{page fault time}]$$
 - **Page fault time** includes
- page fault overhead - swap page out
- swap page in - restart overhead
- Page Fault Time = page fault overhead
= swap page out + swap page in + restart overhead

12. Demand Paging Example

- Memory access time = 200 ns (10^{-9} second)
- Page fault time = 8 millisecond (10^{-3} second)
- Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times (200 \text{ nanosec}) + p \times (8 \text{ millisc})$$

$$= (1 - p) \times (200) + p \times (8,000,000)$$

$$= 200 + 7,999,800 \times p \text{ (in nanosec)}$$


- EAT is directly **proportional** to the **page fault ratio** (p)

13. Page Replacement

- To increase the **degree of multiprogramming**, we are **over-allocating memory**
 - $\Sigma(\# \text{ of pages in each process}) > \# \text{ of frames in the physical memory}$
- In demand-paging system, only **active pages** reside in the physical memory
 - A page is postponed to be loaded until it is used
- When **a new page is loaded** and **there is no available frames**
 - One of pages in the memory needs to be replaced with the new one
 - => Page Replacement
- Page Replacement completes separation between logical memory and physical memory
 - large virtual memory can be provided on a smaller physical memory

14. Need For Page Replacement


- Physical memory
 - Total 8 frames
 - 2 of them for OS
 - 6 are used for users
- Two user processes
 - With 4 pages each
 - 3 of them loaded
- **All frames are used**
- PC indicates an instruction:
 - Load M
 - Page M is on the disk
- Page Replacement needed



15. Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a **page replacement algorithm** to select a **victim frame**
3. Read the desired page into the (newly) free frame and **update** the page and frame tables
4. **Restart** the process

16. Page Replacement procedure



- After finding the target page on the disk and the victim page using page replacement algorithm

1. Swap out the victim page
2. Unset the valid-invalid bit of the victim page
3. Swap the desired page in
4. Set the valid-invalid bit of the target page


17. Page Replacement Algorithms

- Many page replacement algorithm
 - FIFO page replacement
 - Optimal page replacement
 - LRU page replacement
 - LRU-approximation page replacement
 - Additional-Reference-Bits algorithm
 - Second-Chance algorithm
 - Counting-based page replacement
- Performance metric for page replacement algorithm
 - The page fault rate
 - The lower the page-fault rate, the better the performance

- Evaluate algorithm by
 - running it on a particular string of memory references (reference string) and
 - computing the number of page faults on that string
- An example of a reference string is [1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5]

18. Page Faults vs. # of Frames

- 프레임 수 많아지만 frame fault 덜 일어난다



19. 알고리즘들 생략

해보기

30. Global vs. Local Allocation

- Global replacement – process selects a replacement frame from the set of all frames
 - one process can take a frame from another
- Local replacement – each process selects from only its own set of allocated frames

31. Thrashing


- If a process **does not have “enough” frames**, the pagefault rate is very high. This leads to:
 - CPU utilization가 너무 낮으면
 - os가 새로운 프로세스를 시스템에 추가해서 멀티프로그래밍의 degree를 높인다.

- Thrashing : **a process is busy in swapping pages in and out**
 - High paging activity
 - A process is thrashing **if it is spending more time paging than executing**

페이지 폴트가 너무 많이 발생, 실행보다 페이지징에 시간 더걸림

32. Thrashing (Cont.)

- CPU utilization vs. degree of multiprogramming (DM)
 - As DM increases, CPU utilization also increases, until the maximum is reached.
 - As DM is increased even further, thrashing sets in, and CPU utilization drops sharply.
- Decreasing DM can be a solution for the thrashing



33. Demand Paging and Thrashing

- To prevent thrashing, we must provide a process a proper number of frames in demand paging scheme. (프로세스에게 필요로 하는 많은 프레임을 제공해야함)
 - How do we know how many frames it need?
 - Working-Set model with locality is a solution.

- Locality model
 - As a process executes, it moves from locality to locality
 - A locality is **a set of pages that are actively used together**.
 - A program is generally composed of **several different localities**
 - The localities may **overlap**
- Why does thrashing occur?

$$\sum \text{size of locality} > \text{total memory size}$$

34. Locality in a Memory-Reference Pattern

- At a certain period of execution time.
 - Only a set of pages are actively referenced
 - Others are not accessed
- The active pages (**locality**) migrates from one to another as time goes on.
- A program has several different localities
- Localities may **overlap**

35. Page-Fault Frequency Scheme

- Use the **page-fault frequency** (PFF) to **determine the proper number of frames for a process**
 - Establish “**acceptable**” **page-fault rate**
 - If page-fault rate too low, process loses frame
 - If page-fault rate too high, process gains frame
- acceptable rate**

36. Other Issues - Prepaging

- to reduce **the large number of page faults** that occurs at process startup
- prepage **all or some of the pages** a process will need, **before they are referenced**
- But if prepagged pages are unused, **I/O and memory is wasted**
- Assume s pages are prepagged and α ($0 \leq \alpha \leq 1$) of the pages is used
- Is the cost of $s * \alpha$ saving page faults > or < than the cost of prepagging $s * (1 - \alpha)$ unnecessary pages?
- α close to 0 \rightarrow prepagging loses
- α close to 1 \rightarrow prepagging wins

37. Other Issues – Page Size

- Page size selection must take into consideration:
 - **Fragmentation (internal)**
 - page size 사이즈가 크면 내부 단편화 발생
 - **Page table size**
 - page size 줄이면 page 수가 늘어남
 - **I/O overhead**
 - page size가 크면 I/O transfer time이 길어짐
 - **Locality**
 - page size가 크면 locality를 만족시키는 페이지 수가 줄어들음

38. Other Issues – TLB Reach

- **TLB Reach** - **The amount of memory accessible from the TLB**
 - $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
 - Ideally, **the working set of each process is stored in the TLB**
Otherwise there is a high degree of TLB miss
- How to increase the TLB Reach
 - Increase the Page Size.
 - Increase the TLB size
- **Increase the Page Size**
 - This may lead to **an increase in fragmentation** as not all applications require a large page size
- **Provide Multiple Page Sizes**
 - 8KB, 64 KB, 512 KB, 4 MB in Sun UltraSPARC OS
 - This allows **applications that require larger page sizes** the opportunity to use them without an increase in fragmentation

39. Other Issues – I/O interlock

- Pages must sometimes be locked into memory
- Consider I/O - **I/O Interlock**
 - Pages that are used for **copying a file** from a device must be locked from being selected for eviction by a page replacement algorithm.
- **Lock bit**
 - Associated with each page
 - Set to indicate that **the page is locked**
 - Can avoid to be selected by page replacement algorithm